

Netwerken en Systeembeveiliging

Lab 8: Wireless Sensor Network

Chariklis Pittaras (c.pittaras@uva.nl)

Karel van der Veldt (karel.vd.veldt@uva.nl)

Lab date: 27 november, 2012

Hand-in time (return via Blackboard) by 10 december, 2012 23:59CEST

Total points: 40 pts

Abstract

This assignment focuses on peer-to-peer distributed systems using UDP socket programming.

You will learn how to create a network of communicating nodes by simulating a wireless sensor network.

This is a big assignment so you must form a group of two people. On the Blackboard wiki there is a page where groups must enter their names.

Preparation

For this assignment you must use Python. It is already installed on the lab computers, otherwise you can download it from <http://www.python.org>. If you do not know Python, learn it. It is a very simple language.

You can find examples on socket programming in Python in your textbook, section 2.6.

Socket module documentation: <http://docs.python.org/library/socket.html>

Select module documentation: <http://docs.python.org/library/select.html>

Assignment

Wireless sensor networks consist of spatially distributed autonomous sensors that monitor physical conditions such as temperature, light intensity, etc. Your assignment is to program a peer-to-peer program that simulates a sensor node in a distributed system. Every node is identified by its position on a **100x100** grid. All communication will be done using UDP sockets. Every node has a *sensor value*, which is a random number.

Some skeleton code with essential definitions is given in **sensor.py** and **lab8-skeleton.py**. Since every running instance of the program simulates only one node, you have to run multiple instances at the same time to test your code.

Protocol

Nodes communicate by sending messages. The message format is predetermined and has the following fields:

Type	Sequence	Initiator	Neighbor	Operation	Payload
int	int	(int, int)	(int, int)	int	float

- **Type:** The message type (ping, pong, echo, echo_reply).
- **Sequence:** The initiator's echo wave sequence number (task 2).
- **Initiator:** The (x,y) position of the node that initiated a ping (task 1) or echo wave (task 2).
- **Neighbor:** The (x,y) position of a non-initiator node (task 1).
- **Operation:** Echo messages carry an operation for nodes to execute (task 3 and 4).
- **Payload:** Data related to the operation (task 3 and 4).

The file **sensor.py** contains definitions for message length, message types, operation types, and two functions *message_encode()* and *message_decode()* to encode and decode messages to/from a binary format.

All tasks make use of the message format.

GUI Interface

You must use the GUI interface provided by **gui.py** because of the asynchronous nature of the assignment. If you have trouble getting the GUI to work, do not wait until the last moment to let us know.

When the program starts, it must print its own IP:port, position, and sensor value.

In the GUI you must support the following commands:

Command	Action	See task...
ping	Sends a multicast ping message.	1
list	Lists all known neighbors (those who responded to the ping), with (x,y) position and IP:port address.	1
move	Moves the node by choosing a new position randomly.	1

Command	Action	See task...
echo	Initiates an echo wave. All nodes must print to the GUI what messages they receive and the initiator must print if there is a decide event (including the payloads, that will help you with tasks 3 and 4).	2
size	Computes the size of the network.	3
value	The node chooses a new random sensor value (also print it to the GUI).	4
sum	Computes the sum of all sensor values.	4
min	Computes the minimum of all sensor values.	4
max	Computes the maximum of all sensor values.	4

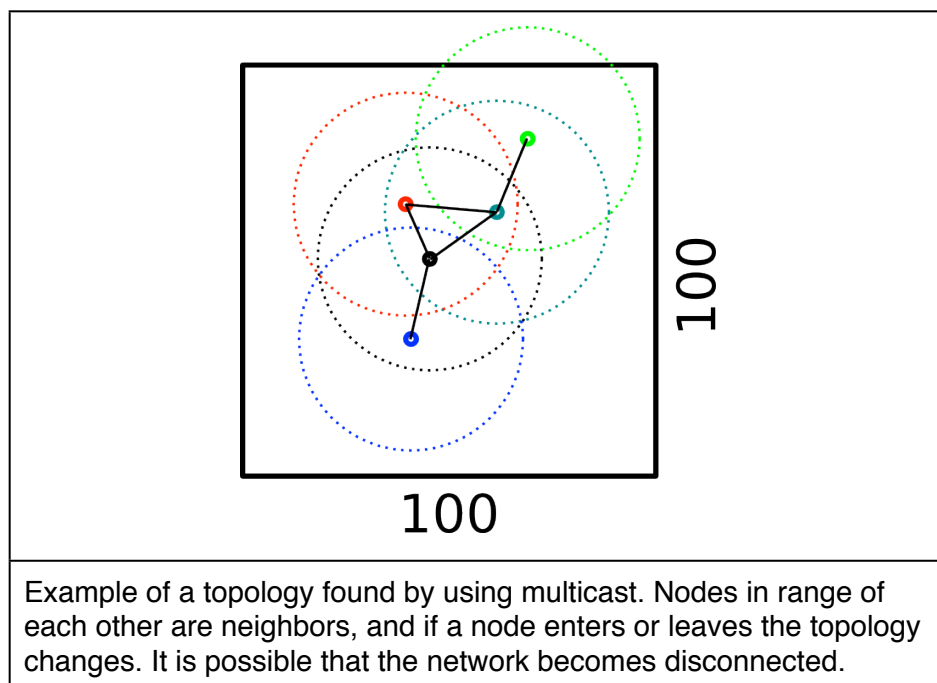
How we will grade

We will run your implementation against our reference implementation.

Therefore it is very important that you implement the protocols and commands exactly as specified.

Task 1 – Neighbor discovery (10 pts)

To communicate, a node must first discover its neighbors (other nodes that are in range). In this assignment we will do this using UDP multicasting. Every node chooses a random position on a **100x100** grid, and has a wireless radius of **50**.



You will need two UDP sockets: one for receiving multicast messages (multicast listener socket), and one for sending multicast messages and sending and receiving unicast messages (peer socket). Setting a UDP socket up for multicast is somewhat obscure, so we provide some code for it in **lab8-skeleton.py**.

The neighbor discovery algorithm works as follows:

1. The initiator clears its list of neighbors and sends a multicast PING message (which contains the initiator's position).
2. When a non-initiator receives the PING message it compares its position with that of the initiator, and sends back a unicast PONG message (which contains the non-initiator's position) to the initiator if they are in sensor range.
3. When the initiator receives a PONG message it adds the position and remote IP:port address to its list of neighbors.

All nodes must periodically (every 5 seconds) resend a PING message to update its list of neighbors, in case some are added or removed. This way the topology updates automatically.

Notes and tips:

- Use `select()` to handle sockets, and set the timeout to zero so that the GUI remains responsive.
- UDP sockets use the `sendto()` and `recvfrom()` methods instead of `send()` and `recv()`.
- Note that PING messages are the *only* messages in the whole assignment that are sent using multicast.
- When in the lab, every group must choose a different multicast port so that different groups do not receive each other's messages. You can use port `MCAST_PORT + group number` as your port number.

Task 2 - Echo algorithm (15 pts)

The echo algorithm is a **centralized wave algorithm** where a message is sent from an initiator and forwarded to all nodes in a distributed system. We will use this algorithm to communicate and perform computations over the entire distributed system.

The echo algorithm works as follows:

1. The *initiator* sends an unicast ECHO message to each of its *neighbors*.
2. When a *non-initiator* receives an ECHO message for the first time, it makes the sender its *father* (for this particular wave). Then it sends an ECHO message to all neighbors except its father. This forwards the wave to the rest of the network.
3. When a non-initiator receives an ECHO message for the first time and has only one neighbor (necessarily the father), it immediately sends an ECHO_REPLY message to the father.

4. When a non-initiator receives an ECHO message from the same wave again, it immediately sends an ECHO_REPLY message to the sender. This way the wave is not propagated twice.
5. When a non-initiator has received an ECHO_REPLY message from all neighbors, it sends an ECHO_REPLY message to its father.
6. When the initiator has received an ECHO_REPLY message from all neighbors, it decides (the algorithm terminates).

The operation is OP_NOOP and the payload is 0 for all messages. The initiator and sequence number must be carried unchanged by all messages, because that is how a wave is uniquely identified.

Echo waves are identified by the tuple (initiator, sequence), which is enough information to handle multiple concurrent waves. Every time a node initiates an echo wave, it increments its sequence number by one.

Note that if even *one* node does not participate in the wave, there will be no decide event (i.e. the algorithm never terminates)! But due to the way you will implement the algorithm, this does not matter: If there is no decide event a node can simply initiate another wave.

Task 3 - Determining the size of the network (10 pts)

The next task is to discover the size of the network using the echo algorithm you have implemented in the previous task. For this you need the message fields **operation** and **payload**. Operation tells nodes which computation to perform, and payload carries data related to that computation. In this task, operation=OP_SIZE (defined in **sensor.py**).

To compute the size of the network using the echo algorithm, every node computes the partial sum of the replies it has received from its neighbors, and sends a message to its parent with payload=1 + sum(payloads received from its neighbors). Eventually the initiator receives the sum of the entire network.

In more detail, the echo algorithm is adapted as follows:

1. The initiator sends an echo with operation=OP_SIZE.
2. When a non-initiator receives a message for the first time, it makes the sender its father for this wave. Then it sends a reply to all neighbors except its father. (this step is unchanged).
3. When a non-initiator receives a message with operation=OP_SIZE for the first time and it has only one neighbor (the father), it immediately sends a reply with operation=OP_SIZE and payload=1.
4. When a non-initiator receives the same message again, it replies with operation=OP_NOOP. Receivers must ignore the payload of messages with operation=OP_NOOP to ensure every node is counted exactly once. Note that a reply must still be sent, or the algorithm never terminates.

5. When a non-initiator has received a reply from all neighbors, it sends a reply to its father with operation=OP_SIZE and payload = $1 + \text{sum}(\text{payloads received from neighbors})$.
6. When the initiator has received a reply from all its neighbors, it computes the size of the network = $1 + \text{sum}(\text{payloads received from all neighbors})$ and decides.

Task 4 - Computations on the network (5 pts)

The final task is to perform computations on the sensor values of each node, in the same way the size of the network is computed in Task 3. Each node randomly generates a sensor value, which is used in the following operations:

Sum (OP_SUM) - Compute the sum of all sensor values.

Minimum (OP_MIN) - Compute the minimum of all sensor values.

Maximum (OP_MAX) - Compute the maximum of all sensor values.

Submission

Submit your working code in a file called **lab8-<yourgroupnumber>.py**.

You must also add your full names and group number at the top of the file.